

第十五讲：死锁和并发错误检测

第 5 节：并发错误检测

向勇、陈渝

清华大学计算机系

xyong,yuchen@tsinghua.edu.cn


2020 年 4 月 12 日

- 1 第 5 节：并发错误检测
 - Concurrency Bug
 - Concurrency Bug Detection
 - AVIO
 - ConSeq & ConMem

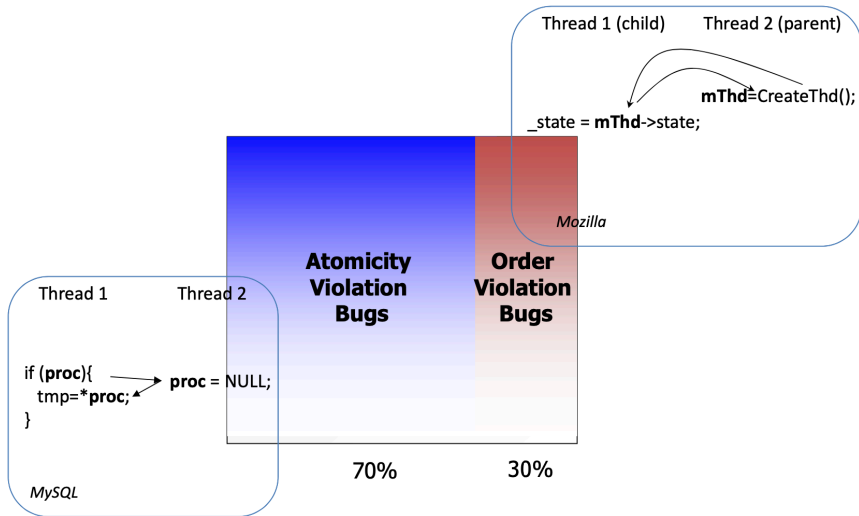
Ref: Shan Lu, Detecting and Fixing Concurrency Bugs, University of Chicago

Concurrency bug

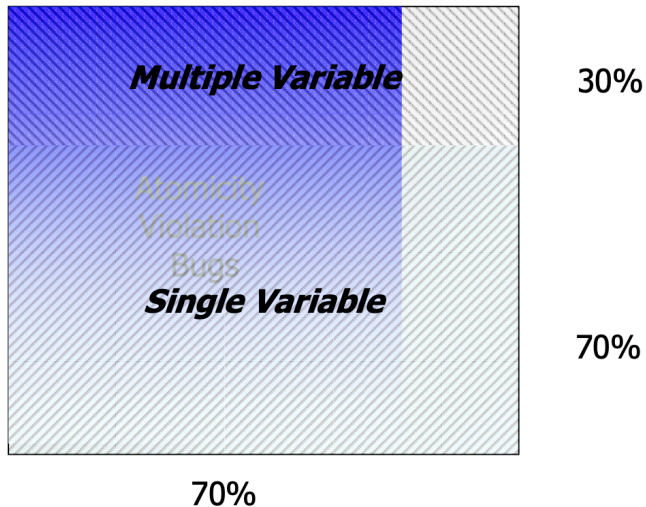
What ordering is guaranteed?



Concurrency bug: Violation



Concurrency bug: Variable



Atomicity Violations

Thread 1

```
if (proc){  
  tmp=*proc;  
}
```

MySQL

Thread 2

```
proc = NULL;
```

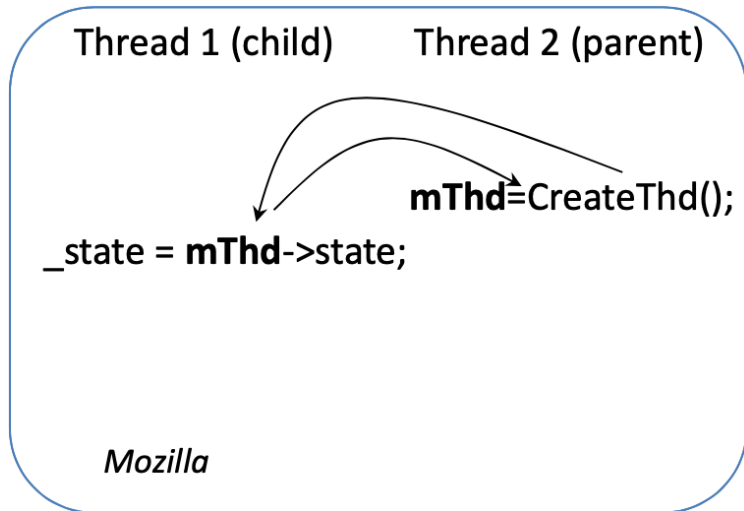
Thread 1

```
while (!flag) {};
```

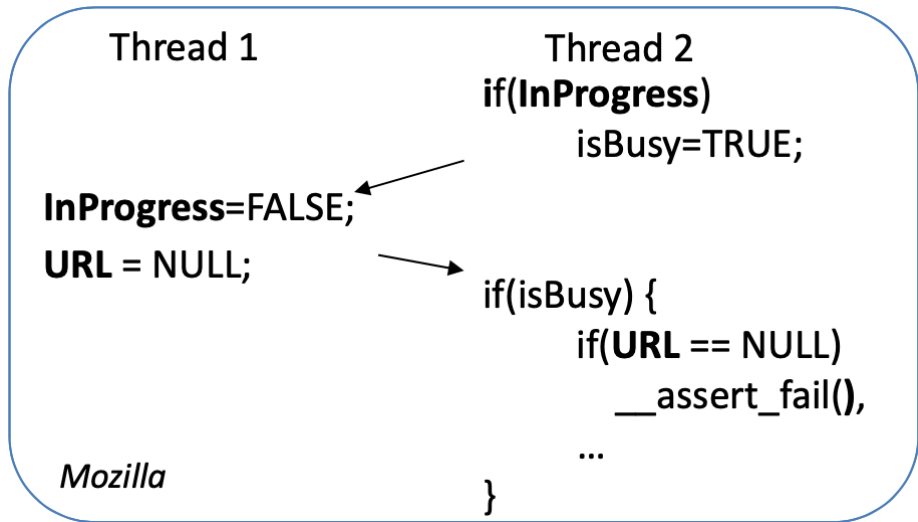
Thread 2

```
flag=TRUE;
```

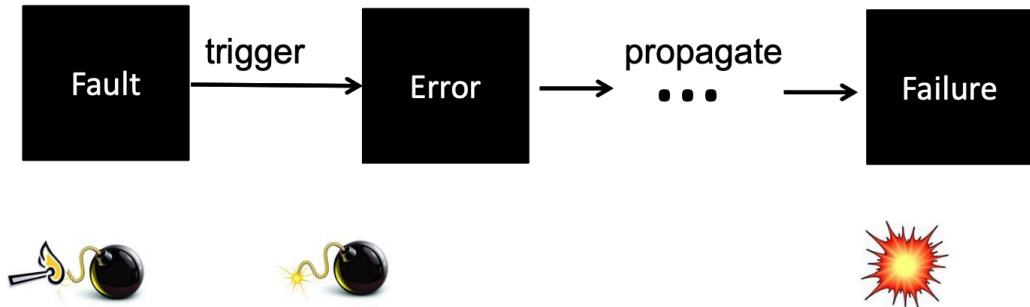
Order Violations



Multi-var Order Violations



The lifecycle of bugs



logical clock algorithm

Use logic time-stamps to find concurrent accesses

Thread 1

Thread 2

lock (L); <0,1>

ptr=NULL; <0,2>

unlock(L); <0,3>

<1,0> **ptr** = malloc(10);

<2,3> lock (L);

<3,3> **ptr**[0]='a';

<4,3> unlock(L);

Lock-set algorithm

A common lock should protect all conflicting accesses to a shared variable

Thread 1

Thread 2

lock (L);

ptr=NULL; <L>

unlock(L);

</> **ptr** = malloc(10);

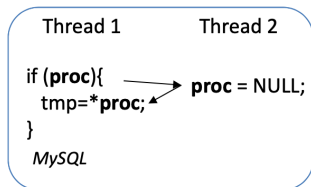
lock (L);

<L> **ptr**[0]='a';

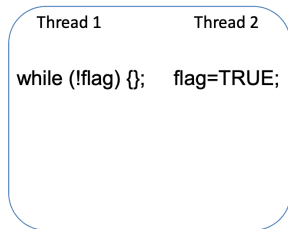
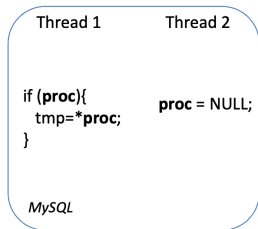
unlock(L);

How to detect atomicity-violations?

Know which code region should maintain atomicity



Judge whether a code region's atomicity is violated



AVIO: Detecting Atomicity Violations via Access-Interleaving Invariants (ASPLOS' 06)

Atomicity violation = unserializable interleaving



AVIO: Detecting Atomicity Violations via Access-Interleaving Invariants (ASPLOS' 06)

Totally 8 cases of interleaving

Read x
Read x
Read x

Write x
Read x
Read x

Read x
Write x
Read x

Write x
Write x
Read x

Read x
Read x
Write x

Write x
Read x
Write x

Read x
Write x
Write x

Write x
Write x
Write x

AVIO: Detecting Atomicity Violations via Access-Interleaving Invariants (ASPLOS' 06)

4 out of 8 cases are interleaving violations

Read x
Write x
Read x

Inconsistent
views

Write x
Write x
Read x

Too early
overwritten

Write x
Read x
Write x

Leaking
intermediate value

Read x
Write x
Write x

Using stale
value



Both hardware and software solutions exist

If we cannot find a more accurate **root-cause** pattern, let's look at the **effect** patterns of concurrency bugs!

- ConMem
 - Detecting Severe Concurrency Bugs through an Effect-Oriented Approach, ASPLOS' 10
- ConSeq
 - Detecting Concurrency Bugs through Sequential Errors, ASPLOS' 11

The lifecycle of concurrency bugs: Fault

based on 70 real-world bugs





Data races
Atomicity violations
 single variable
 multiple variables
Order violations

...

The lifecycle of concurrency bugs: Error



based on 70 real-world bugs



- Memory errors
 - NULL ptr
 - Dangling ptr
 - Uninitialized read
 - Buffer overflow
- Semantic errors

The lifecycle of concurrency bugs: Failure

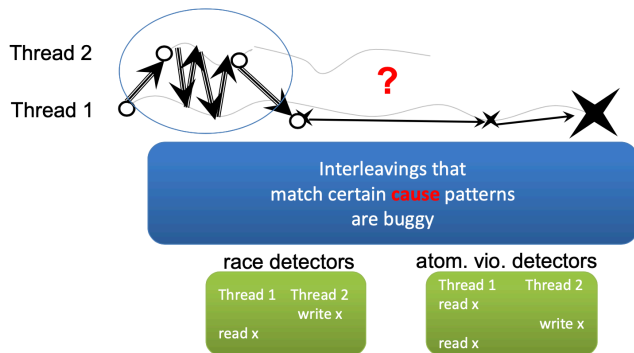
based on 70 real-world bugs



*short
single-threaded*

- ☀ Crash @ invalid memory
- ☀ Crash @ assertion
- ☀ Infinite loops
- ☀ Incorrect outputs
- ☀ Error messages


Cause-oriented approach



Limitations

- False positives
- False negatives


Effect-oriented approach



- Step 1: Statically identify potential failure/error site
- Step 2: Statically look for critical reads
- Step 3: Dynamically identify buggy interleaving

ConMem

Detecting Severe Concurrency Bugs through an Effect-Oriented Approach, ASPLOS' 10




- Memory errors
 - NULL ptr
 - Dangling ptr
 - Uninitialized read
 - Buffer overflow

- Semantic errors

- Crash @ invalid memory
- Crash @ assertion
- Infinite loops
- Incorrect outputs
- Error messages

ConSeq: Detecting Concurrency Bugs through Sequential Errors, ASPLOS' 11



- Memory errors
 - NULL ptr
 - Dangling ptr
 - Uninitialized read
 - Buffer overflow

- Semantic errors

- Crash @ invalid memory
- Crash @ assertion
- Infinite loops
- Incorrect outputs
- Error messages