

第二十一讲：异步编程 (Asynchronous Programming)

第 1 节：Background

向勇、陈渝

清华大学计算机系

xyong,yuchen@tsinghua.edu.cn

2020 年 5 月 6 日

提纲

- 1 Background
- 2 Futures in Rust
- 3 Generators and `async/await`
- 4 Self-Referential Structs & `Pin`
- 5 Waker and Reactor

Ref:

- [Futures Explained in 200 Lines of Rust](#), by Carl Fredrik Samson
- [Writing an OS in Rust - Async/Await](#), by Philipp Oppermann
- [Zero-cost futures in Rust](#), by Aaron Turon
- [Rust's Journey to Async/Await](#), by Steve Klabnik
- [Asynchronous Programming in Rust](#)

recap: Multitasking

Non-Preemptive multitasking

- The programmer 'yielded' control to the OS
- Every bug could halt the entire system
- Example: Windows 95

Preemptive multitasking

- OS can stop the execution of a process, do something else, and switch back
- OS is responsible for scheduling tasks
- Example: UNIX, Linux

recap: User-level Thread

Advantages

- Simple to use
- A "context switch" is reasonably fast
- Each stack only gets a little memory
 - You can have hundreds of thousands of user-level threads running
- Easy to incorporate preemption

Drawbacks

- The stacks might need to grow
 - Solving this is not easy and will have a cost
- Need to save all the CPU state on every switch
- Complicated to implement correctly if you want to support many different platforms

Example: **Green Threads**

recap: Kernel-supported Threads

Advantages

- Easy to use
- Switching between tasks is reasonably fast
- Getting parallelism for free

Drawbacks




- OS level threads come with a rather large stack
- There are a lot of syscalls involved
- Might not be an option on some systems, such as http server

Example:

- Using OS threads in Rust

recap: What is async?

- Parallel: do multiple things at once
- Concurrent: do multiple things, not at once
- Asynchronous: Describe lang/prog features that enable parallelism & concurrency
- Task: Some computation running in a parallel or concurrent system



Why we need async?



C10K Problem in 1999 ... C100K, C1M, C10M, C100M ...

- 网络服务在处理数以万计的客户端连接时，往往出现效率低下甚至完全瘫痪，这被称为 C10K 问题
- C10K 问题的提出者 Dan Kegel：软件工程师
- Web1.0 Ok! Web2.0 Cry!

Why we need async?

C10K Problem in 1999 ... C100K, C1M, C10M, C100M ...


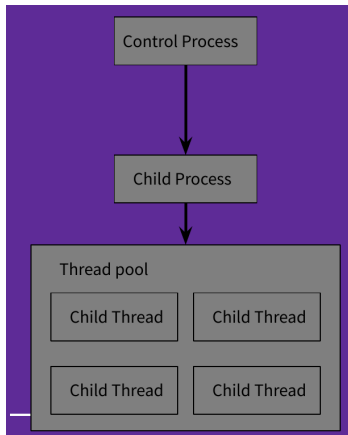
- 网络服务在处理数以万计的客户端连接时，往往出现效率低下甚至完全瘫痪，这被称为 C10K 问题
- C10K 问题的提出者 Dan Kegel：软件工程师
- Web1.0 Ok! Web2.0 Cry! 核心问题：**时间开销** + **空间开销**



Why we need async?

解决方法: C10K Problem in 1999 ... C100K, C1M, C10M, C100M ...


- Serve one client with each thread/process, and use blocking I/O : Apache、ftpd



Why we need async?

解决方法: C10K Problem in 1999 ... C100K, C1M, C10M, C100M ...

- Serve many clients with each thread, and use asynchronous I/O : nginx



- Event-driven
- Asynchronous
- Non-blocking

Why we need async?

解决方法: C10K Problem in 1999 ... C100K, C1M, C10M, C100M ...

- Serve many clients with each thread, and use asynchronous I/O : nginx



Who provide async mechanism?

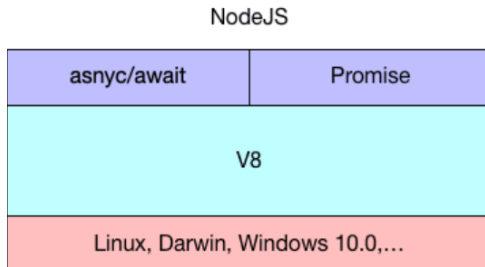
Syntax	Type
Runtime + Thread Pool	
Kernel Abstractions	
Kernel	

- Implement your own way of handling threads and queues on program level (green threads)
- Add syntactic sugar to your language so the runtime/compiler can identify async parts of the code
- Add async types so they can notify when they are "done"

F# added to the core design in 2007: computation expressions and their application to asynchronous programming

Who provide async mechanism?


Syntax	Type
Runtime + Thread Pool	
Kernel Abstractions	
Kernel	



```
const async_method = async () => {  
  const dbResults = await dbQuery();  
  const results = await serviceCall(dbResults);  
  console.log(results);  
}
```

Who provide async mechanism?

Syntax	Type
Runtime + Thread Pool	
Kernel Abstractions	
Kernel	



```
f(greeting string) {  
    fmt.Println(greeting, ", World!")  
}  
  
go f("Hello")
```

Who provide async mechanism?

Syntax	Type
Runtime + Thread Pool	
Kernel Abstractions	
Kernel	

Rust

async/await	Future
tokio, romio + juliex	
mio	
Linux, Darwin, Windows 10.0,...	

```
async fn hello_world() {  
    let x: u8 = foo().await;  
    println!("{}", hello, world!",x);  
}  
  
fn main() {  
    let future = hello_world(); // do nothing  
    block_on(future); // print something  
}
```

Async Prog: Callback based approaches

A callback based approach is to save a pointer to a set of instructions we want to run later together with whatever state is needed.

Advantages

- Easy to implement in most languages
- No context switching
- Relatively low memory overhead

Drawbacks

- Memory usage grows linearly with the number of callbacks
 - Each task must save the state it needs for later
- Callback hell: Hard to debug
- Require a substantial rewrite to go from a "normal" program flow to one that uses a "callback based" flow

Example: **Callback based approaches**

From callbacks to futures (deferred computation)

A callback based approach.

```
\\JavaScript
setTimeout(200, () => {
  setTimeout(100, () => {
    setTimeout(50, () => {
      console.log("I'm the last one");
    });
  });
});
```

From callbacks to futures (deferred computation)

Promises: deal with the complexity which comes with a callback based approach.

```
\\JavaScript
function timer(ms) {
  return new Promise((resolve) => setTimeout(resolve, ms));
}

timer(200)
  .then(() => return timer(100))
  .then(() => return timer(50))
  .then(() => console.log("I'm the last one"));
```

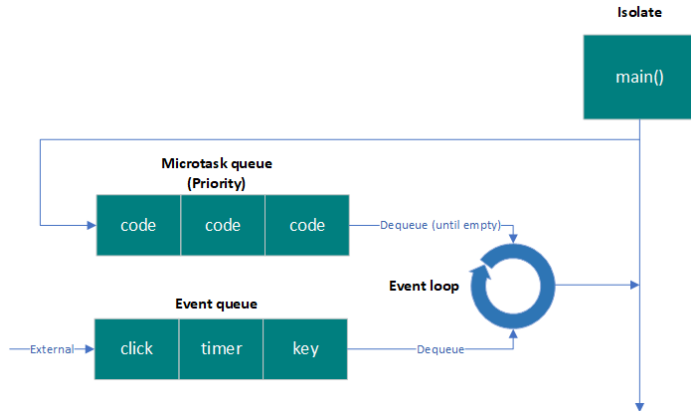
From callbacks to futures (deferred computation)

Promises: deal with the complexity which comes with a callback based approach.

```
\\javascript
  async function run() {
    await timer(200);
    await timer(100);
    await timer(50);
    console.log("I'm the last one");
  }
```

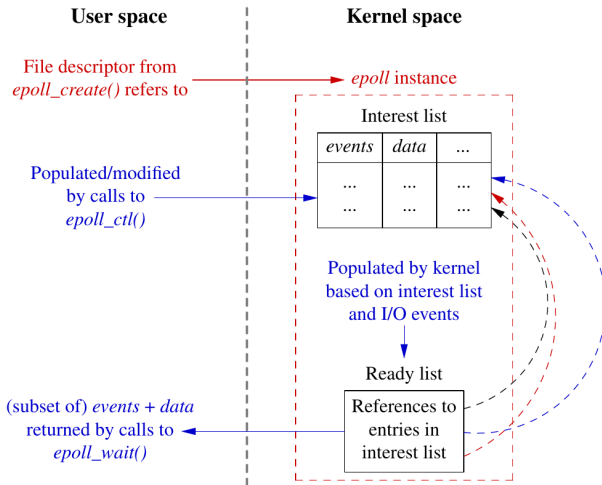
- The 'run' function as a *pausable* task consisting of several sub-tasks
 - On each "await" point it yields control to the scheduler
- When the sub-tasks changes state to either 'fulfilled' or 'rejected', the task is scheduled to continue to the next step

Event queue: Epoll, Kqueue and IOCP



- Epoll is the Linux way of implementing an event queue
- Kqueue is the MacOS way of implementing an event queue
- IOCP or Input Output Completion Ports is the way Window handles event queue

Epoll



Procedure for read data from a socket using epoll

- 1 Create an event queue by calling the syscall 'epoll_create' or 'kqueue'
- 2 Ask the OS for a file descriptor representing a network socket
- 3 Register an interest in 'Read' events on this socket
 - In order to receive a notification when the event is ready in the event queue we created
- 4 Call 'epoll_wait' or 'kevent' to wait for an event
 - Block (suspend) the thread it's called on
- 5 When the event is ready, our thread is resumed, and return from our "wait" call with data about the event
- 6 Call 'read' on the socket we created

Example

- Epoll example
- Complete example